

Heuristic Search Practice. TSP with A*, SPA* and GA*

Israel Solís Iglesias^[UO282162], Omar Teixeira González^[UO281847], and David Leszek Warzynski Abril^{[UO278968]**}

Intelligent Systems, School of Computer Engineering, University of Oviedo, Oviedo, Spain

{U0282162,U0281847,U0278968}@uniovi.es
<https://ingenieriainformatica.uniovi.es/>

Abstract. The Traveling Salesman Problem (TSP) poses the challenge of a traveler commencing from a city and navigating through all available cities exactly once, ultimately returning to the initial departure city, all while endeavoring to identify the most cost-efficient route. In this article, we will conduct an experimental investigation into solving the TSP utilizing the A* search algorithm and SPA*. These algorithms will be implemented with various heuristics to explore state spaces effectively. Lastly, we will present and compare the results obtained from employing these distinct methods.

Keywords: A* · SPA* · GA · State space search · Heuristic · TSP.

1 Introduction

Search algorithms are the foundation of every intelligent system, categorized into two distinct types: uninformed/blind and informed algorithms. This paper delves into informed algorithms, explicitly focusing on A* and SPA*, employing various heuristics and their application to the Traveling Salesman Problem (TSP). A subsequent comparison between these algorithms will be conducted, alongside an exploration of applying genetic algorithms, providing a comprehensive analysis.

This paper will present comparative analyses through carefully executed experiments, showcasing outcomes in dedicated sections.

The structure of this article unfolds as follows. Section 2 initiates with an elucidation of the traveling salesman problem. Sections 3, 4, and 5 sequentially expound on search algorithms, their application to the TSP, and the findings of the experimental study. Finally, Section 6 encapsulates the conclusions drawn from this research.

2 Traveling Salesman Problem (TSP)

The Traveling Salesman Problem (TSP) is an algorithmic challenge that involves determining the most efficient route, to visit each city exactly once and return to

* Supported by University of Oviedo

** The authors have been arranged in alphabetical order based on their surnames

the starting city. The problem dates back to the 18th century, with notable early contributors including Sir William Rowan Hamilton and Thomas Penyngton. However, it wasn't until the 1930s, that Karl Menger delved into the TSP's general form as it is studied today[6].

Over the years, extensive research has revealed a myriad of applications for the TSP, ranging from the drilling of printed circuit boards to computer wiring and vehicle routing.

The TSP can be tackled using both uninformed algorithms, such as Depth-First, and informed algorithms like A* and SPA*. Additionally, genetic algorithms offer another avenue for solving the problem, as elaborated in the subsequent sections[1].

3 Search algorithms

3.1 A* algorithm

The A* algorithm is a specialization of the Breadth-First (BF) algorithm which incorporates an evaluation heuristic function denoted as $f(n) = g(n) + h(n)$ with:

- $g(n)$. the cost from the start node to n .
- $h(n)$. the estimate of the cost from n to the end node.

Regarding its properties, A* is a complete algorithm, which means, it guarantees to find a solution when there is one. Both its admissibility and consistency are determined by the heuristic $h(n)$ used.

To ensure its admissibility, $h(n)$ must be optimistic, which is to say that $h(n) < h^*(n) \forall n$, where $h^*(n)$ is the best cost from n . On the other hand, its consistency is assured when $h(n) > h(n')$ for any node n' , a child of a node n . When $h(n)$ is admissible, consistent, or both, A* will always yield the optimal solution. Besides, when it is consistent A* will be optimally efficient compared to other similar algorithms with no all "non-pathological" search problems [2].

3.2 SPA* algorithm

Static Ponderation A*[5] is classified among ϵ -admissible algorithms, representing a relaxed variant of the A* algorithm. This adaptation becomes necessary when dealing with large-scale problems, compelling a relaxation of the algorithm's admissibility (while still requiring an admissible heuristic). Essentially, while A* guarantees an optimal solution, it may explore more paths to find it. By relaxing this condition, the algorithm traverses fewer nodes, potentially yielding a solution that is not necessarily optimal.

ϵ -admissible algorithms share the characteristic that the obtained solution will have a maximum cost of $(1 + \epsilon)C^*$, where $\epsilon > 0$. In essence, a larger ϵ results in the expansion of fewer nodes on the path to the solution.

In this paper, the SPA* algorithm used, is defined as $f(n) = g(n) + (1 + \epsilon)h(n)$.

4 Application of algorithms to the TSP

Both search and genetic algorithms have been utilized to discover the path with the minimum cost that traverses all cities and returns to the initial one.

4.1 Search algorithms

A* and SPA* algorithms[3] share the same search space and heuristics. However, there is a unique aspect in the case of SPA*, which will be detailed later in this document.

The algorithm initiates by generating the search space, constructing a graph where each node represents a new addition to the path. Four implemented heuristics determine the path, all designed to identify the path with the minimum cost. These four heuristics are as follows:

- **h1**. Based on calculating the sum of edges connecting the remaining cities to be visited, including the current one.
- **h2**. The sum of the $N - k + 1$ least costly edges in the residual graph.
- **h3**. The sum of the $N - k + 1$ least costly edges in the residual graph, ensuring that each city in the residual graph has a connecting edge.
- **h_mst**. Cost of a minimum spanning tree of the residual graph.

Specifically, the h3 heuristic consists of giving a list of edges composing the residual graph sorted by cost, progressively adding the edges of the remaining cities, and removing these cities as they are added. Once all cities are included, if the total length of the stored edges is less than $N - k + 1$, where N is the total number of cities and k is the number of cities visited, then add the remaining edges with the lowest cost from the residual graph, regardless of the cities they connect.

4.2 Genetic algorithms

The genetic algorithm is segmented into various components, with the initial part focusing on the encoding scheme. This scheme is founded on permutations of the cities $(1, 2, \dots, n - 1)$.

Another crucial element is the crossover operator, with two operators implemented in this case: Order crossover and Uniform crossover, the first one copies a substring of symbols from the first parent to the child, maintaining order and position, the rest occupy the remaining positions keeping the order they have in the second parent, while the uniform crossover takes a x number of values from the first parent and then takes the rest in relative order from the other parent.

The final component is the fitness function, which calculates the cost of the path from the initial city to itself, navigating through the remaining cities in the order specified by the individual.

5 Experimental study

5.1 Experimental study design

In the experimentation of this article, the methods described in the previous sections will be compared. Different instances of the TSP will be solved for this comparison: three of them were provided in the class exercises –gr17, gr21 without four cities and gr21 with all cities–, and the last one –fri26–, obtained from [4], which also had the last three cities removed.

In all the experiments, two metrics will be collected for the comparison: the execution time and the cost of the best solution found by the method. The execution time and the best solution found in the case of genetic algorithms are something that can vary with each execution. Therefore, these two metrics have been obtained from 10 different runs, and in the results to be presented, you will find the mean and standard deviation over the 10 runs. In the heuristic algorithms, the number of expanded nodes will also be collected.

Specifically, the SPA* algorithm has been employed with all four aforementioned heuristics and with five different ϵ values (0.2, 0.4, 0.6, 0.8, 0.99) for each heuristic.

The primary specifications of the computer used for experimentation are as follows:

- Processor AMD Ryzen 5 3600 6-Core Processor, 3600 Mhz, 6 Core(s), 12 Logical Processor(s).
- RAM 16.0 GB.

Talking genetic algorithms, have been used with different parameters, such as the number of population, which has been set to 100, or the number of generations, which oscillates between the values (100, 150, 200, 250, 300). Being more concrete about implementation, it has also been used elitism, to take the best out of the children, mutation (where it changes 2 values from the child) with 10% probability, and the possibility of choosing among the 2 aforementioned crossover operators.

5.2 Experimental results

Heuristic comparison The results unequivocally indicate that *h2* emerges as the least effective heuristic in terms of expanded nodes. Conversely, *h_mst* stands out as the most efficient heuristic regarding expanded nodes. However, this advantage is counterbalanced by the time it takes to reach the solution, which can be longer than that of other heuristics in certain instances, as illustrated in Table 1.

Table 1. A* Performance Across Multiple Instances of Varying Sizes and Heuristics.

Instance	Heuristic	Expanded Nodes	Average (Time)	Standard Deviation (Time)	Length (Cost)
gr17	h1	94216	11,0028	0,72648496	2085
	h2	248137	23,78591528	1,184029394	
	h3	83744	13,20123687	1,505892807	
	h_mst	30413	57,72750638	2,823759739	
gr21-4	h1	3771	0,718152571	0,540446349	2369
	h2	25510	23,78591528	1,184029394	
	h3	3108	1,019539356	0,591314004	
	h_mst	269	1,551998925	0,00507672	
gr21	h1	96186	21,91698604	0,795147873	2707
	h2	561594	91,07996345	4,619105429	
	h3	58685	22,08413286	1,079720095	
	h_mst	7202	66,10514495	0,68714398	
fri26-5	h1	956790	150,8601679	6,84548460	660
	h2	517093	84,11687043	2,454384717	
	h3	20254	9,457087541	0,672234078	
	h_mst	788	9,617299557	0,405164665	

Heuristic resume Upon examining the results in Table 1, we can preliminary assert the admissibility of the heuristics. Concerning monotony, it can be preliminarily affirmed that h1, h2, and h_mst exhibit this property. However, h3, on the other hand, does not achieve monotony due to node re-expansion. Fig. 1.

Epsilon comparison The findings demonstrate that a value of $\epsilon = 0.2$ is the least favorable when considering expanded nodes and time, but it excels in converging toward the optimal solution. In contrast, $\epsilon = 0.99$ stands out as the optimal choice in terms of expanded nodes and time efficiency. However, this advantage is mitigated by the resulting solution's cost, which may significantly deviate from the optimal solution, as depicted in Table 2 and Fig. 1.



**Fig. 1.** Comparison of SPA* Costs with Various ϵ Values to A* Optimal Cost for the GR17 Instance.

Table 2. SPA* Results for GR17 Instance with Various Heuristics and Epsilons. Time measurements are in seconds.

Heuristic	ϵ	Expanded Nodes	Average (Time)	Standard Deviation (Time)	Length (Cost)
h1	0.2	30273	3.545046353	0.223812911	2085
	0.4	11853	1.549099922	0.17057407	2103
	0.6	4518	0.55580008	0.167406067	2238
	0.8	1636	0.227800226	0.128229416	2238
	0.99	478	0.077399087	0.095782597	2268
h2	0.2	165565	16.91656694	1.121160929	2088
	0.4	99602	11.38393157	1.072885589	2088
	0.6	58289	7.193899703	0.580293861	2088
	0.8	37246	4.845499349	0.634438286	2088
	0.99	21205	3.06549921	0.439285928	2120
h3	0.2	37236	6.12194941	0.700358208	2088
	0.4	16477	2.947298694	0.02590811	2088
	0.6	7335	1.345900083	0.034099174	2146
	0.8	1643	0.345799541	0.020394877	2146
	0.99	554	0.119900203	0.018272382	2151
h_mst	0.2	11544	26.34758334	0.114694006	2090
	0.4	4346	10.04292498	0.051799545	2090
	0.6	1949	4.403960323	0.038079103	2090
	0.8	1495	4.094906282	0.033538196	2090
	0.99	730	2.903900433	0.047827413	2218

Algorithm comparison A* vs SPA* As evident from Table 1 and Table 2, the utilization of SPA* significantly reduces the number of expanded nodes and the time required to obtain a solution, although at the expense of results that deviate further from optimal. The disparity in expanded nodes is visually depicted in Fig. 2.

**Fig. 2.** Comparison of Expanded Nodes: SPA* with Various ϵ vs. A* for GR17 Instance Using Four Defined Heuristics.

Number generations genetic comparison The results indicate that with a higher number of generations, the approach to the optimal cost of the problem improves. However, this improvement is reflected in the time required to achieve it. For instance, with 300 generations, the best cost is closer to the optimum compared to 100 generations, but the time taken is greater with 300, as illustrated in Tab 3, Fig. 3, and Fig. 4.

Crossover Operator comparison OX vs UX The results reveal that *OX* is marginally more effective in achieving a solution close to the optimum compared to *UX*. However, the times obtained by both operators across different generations are also relatively comparable. The nuanced difference is apparent in Table 3, Fig. 3, and Fig. 4.

Table 3. Genetic Algorithm Results for GR17 with OX and UX Crossover Operators, Varied Number of Generations.

Crossover Op	N Gens	Best (Avg)	Solution Cost (Avg)	Standard Deviation (Cost)	Average (Time)
OX	100	2497.5	2698.686	167.8344766	1.568090987
	150	2453.3	2681.111333	164.0302266	2.328969046
	200	2391.9	2552.1625	178.4466897	3.087874409
	250	2416.9	2572.3112	171.8446308	3.909895103
	300	2245.5	2448.593	245.3503937	4.703596943
UX	100	2548.5	2824.911	199.4617631	1.475074123
	150	2519.3	2732.484667	196.5803471	2.203592419
	200	2430.5	2631.839	218.8820025	2.941896438
	250	2369.7	2621.2668	228.8854709	3.694080231
	300	2444.5	2659.570333	201.2008214	4.562186526




Fig. 3. Average Cost Comparison between OX and Different Generations for GR17 Instance, with Four Series Based on Maximum Generation Counts, Alongside the Optimal A* Solution Cost.




Fig. 4. Average Cost Comparison between UX and Different Generations for GR17 Instance, with Four Series Based on Maximum Generation Counts, Alongside the Optimal A* Solution Cost.

6 Conclusions

As has been demonstrated, the findings presented in this paper lead to the conclusion that employing SPA* allows for the attainment of solutions that may be less precise but in a shorter time compared to A*. A* with the specified heuristics and instances consistently reaches the optimal solution.

Furthermore, the application of genetic algorithms demonstrates the capability to address complex optimization problems. It's noteworthy that achieving a more precise solution is feasible with an increased number of populations and generations in the genetic algorithm approach.

References

1. Davendra, D.: Traveling salesman problem: Theory and applications (2010)
2. Dechter, R., Pearl, J.: Generalized best-first search strategies and the optimality of a*. J. ACM **32**(3), 505–536 (jul 1985). <https://doi.org/10.1145/3828.3830>
3. Education, S.: Amit's a* pages, <http://theory.stanford.edu/~amitp/GameProgramming>
4. Grant, R.: Tsplib 95 (2022), <https://github.com/rhgrant10/tsplib95>
5. Russell, S.J., Norvig, P.: Artificial intelligence a modern approach (2010)
6. of Waterloo, U.: Tsp, <https://www.math.uwaterloo.ca/tsp/index.html>